

DEEP  
LEARNING  
WORKSHOP

Dublin City University  
27-28 April 2017



#InsightDL2017

Day 2 Lecture 8

# Recurrent Neural Networks



**Xavier Giro-i-Nieto**

[xavier.giro@upc.edu](mailto:xavier.giro@upc.edu)

Associate Professor

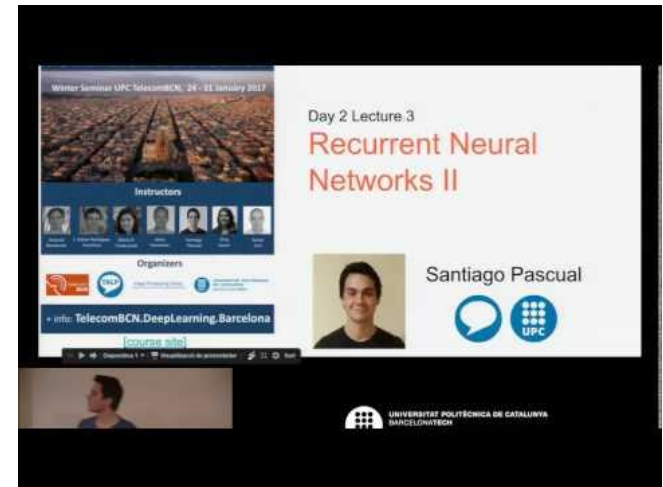
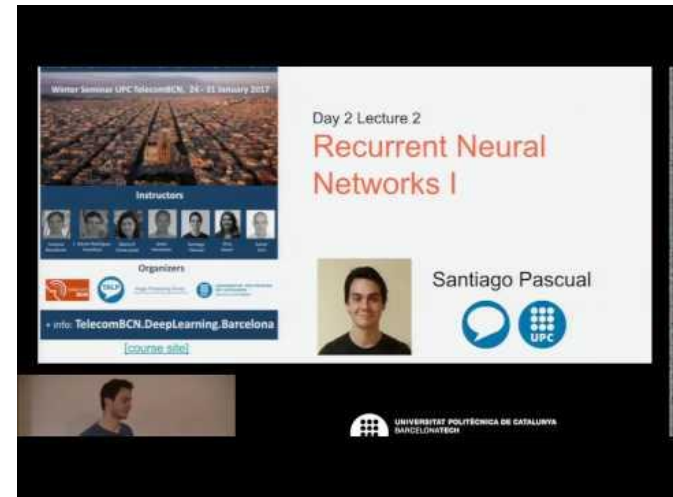

Universitat Politècnica de Catalunya  
Technical University of Catalonia



# Acknowledgments



Santiago Pascual



# Outline

1. The importance of context
2. Where is the memory?
3. Vanilla RNN
4. Problems
5. Gating methodology
  - a. LSTM
  - b. GRU

# The importance of context

- Recall the 5th digit of your phone number
- Sing your favourite song beginning at third sentence
- Recall 10th character of the alphabet

Probably you went straight from the beginning of the stream in each case...  
because in sequences order matters!

Idea: retain the information preserving the importance of order

# Recall from Day 1 Lecture 4...


Every  $\mathbf{y}/\mathbf{h}_i$  is computed from the sequence of forward activations out of input  $\mathbf{x}$ .

$$\mathbf{y} = f(\mathbf{W}_3 \cdot \mathbf{h}_2 + \mathbf{b}_3)$$

$$\mathbf{h}_2 = f(\mathbf{W}_2 \cdot \mathbf{h}_1 + \mathbf{b}_2)$$

$$\mathbf{h}_1 = f(\mathbf{W}_1 \cdot \mathbf{x} + \mathbf{b}_1)$$


Feed FORWARD



Slide Credit: Hugo Laroché NN course


# Where is the Memory?

If we have a sequence of samples...




predict sample  $x[t+1]$  knowing previous values  $\{x[t], x[t-1], x[t-2], \dots, x[t-\tau]\}$

# Where is the Memory?




Feed Forward approach:

- static window of size  $L$
- slide the window time-step wise




# Where is the Memory?



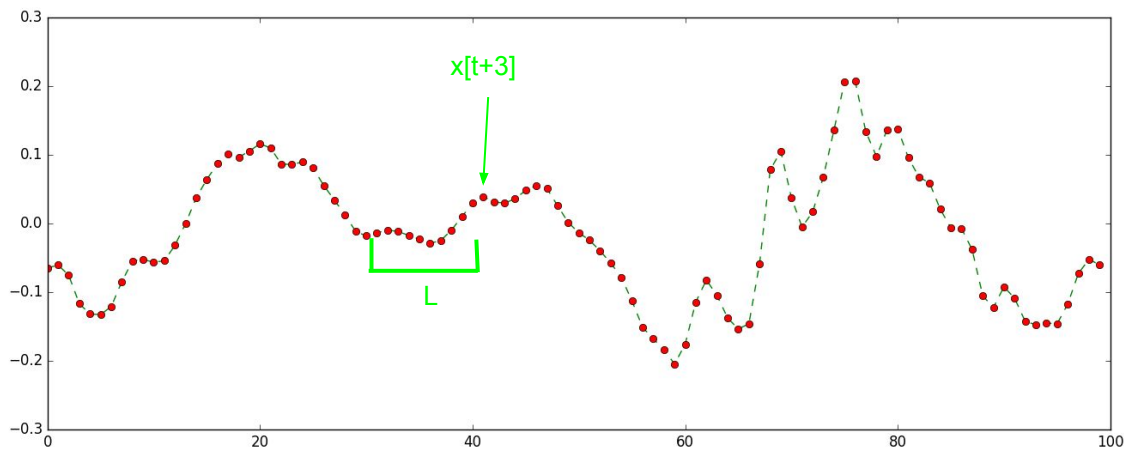
Feed Forward approach:

- static window of size  $L$
- slide the window time-step wise






# Where is the Memory?



Feed Forward approach:


- static window of size  $L$
- slide the window time-step wise



# Where is the Memory?

Problems for the feed forward + static window approach:

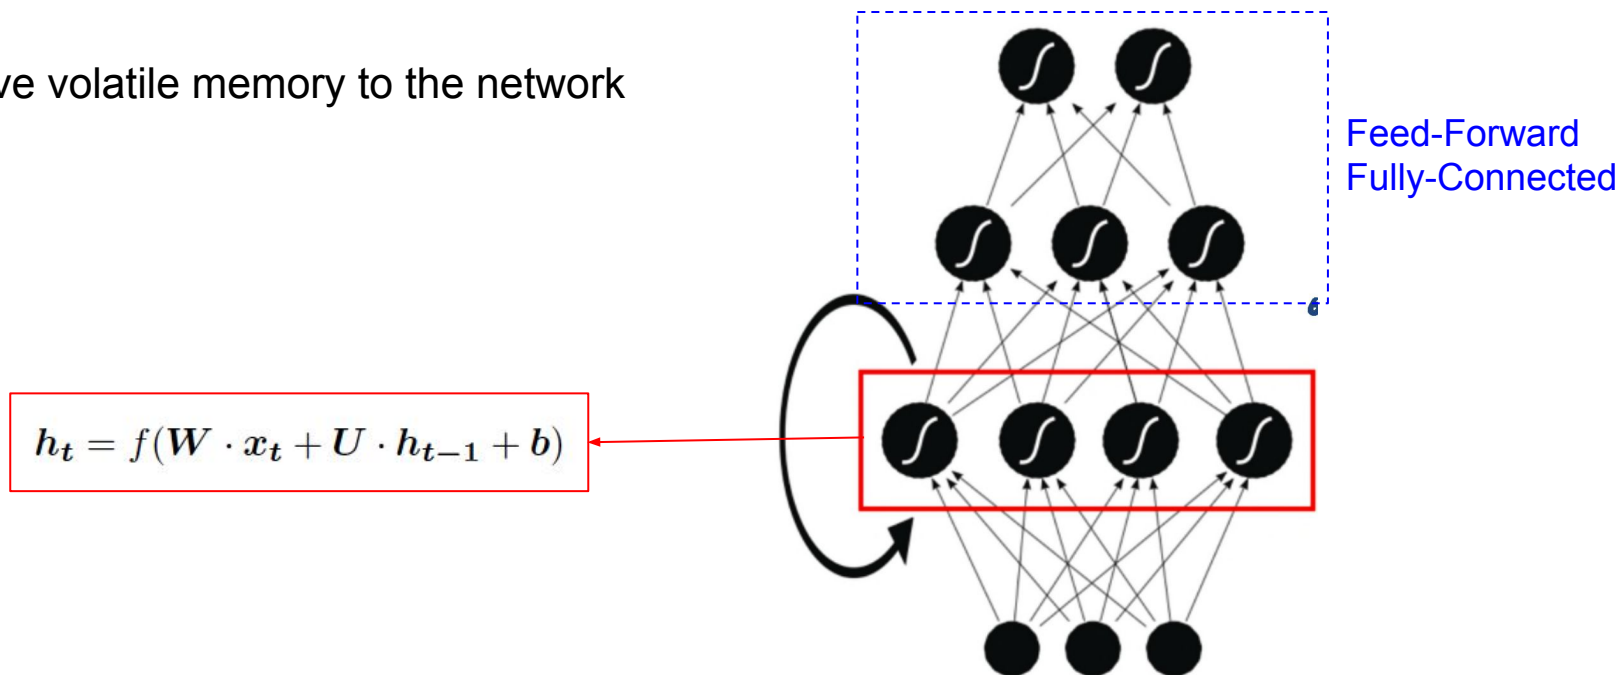
- What's the matter increasing  $L$ ? → Fast growth of num of parameters!
- Decisions are independent between time-steps!
  - The network doesn't care about what happened at previous time-step, only present window matters → doesn't look good
- Cumbersome padding when there are not enough samples to fill  $L$  size
  - Can't work with variable sequence lengths



# Recurrent Neural Network

Solution: Build specific connections capturing the temporal evolution → **Shared weights in time**

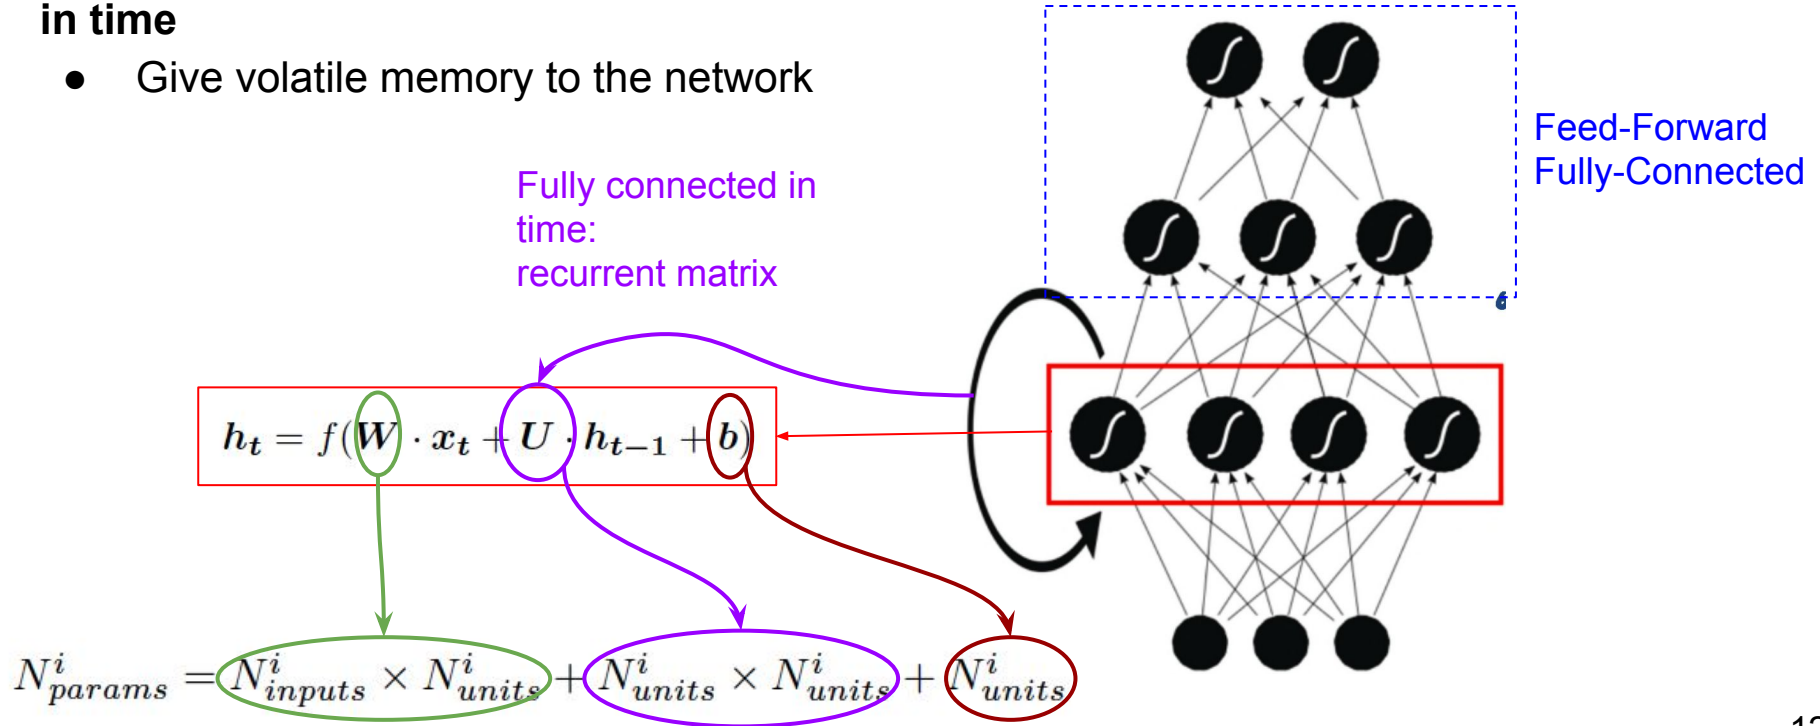
- Give volatile memory to the network



# Recurrent Neural Network



Solution: Build specific connections capturing the temporal evolution → **Shared weights in time**

- Give volatile memory to the network




# Recurrent Neural Network

Front View



time

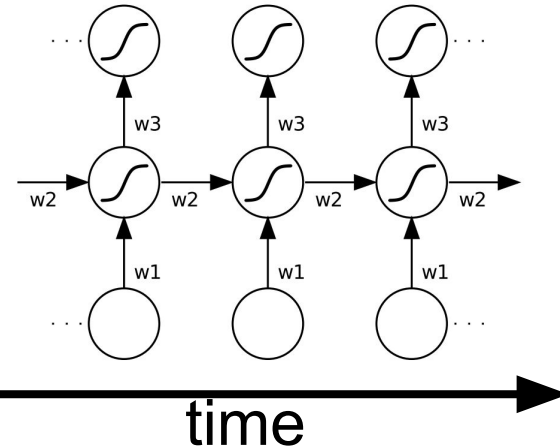


Slide credit: Xavi Giro

Rotation  
 $90^\circ$

Rotation  
 $90^\circ$


Side View



# Recurrent Neural Network

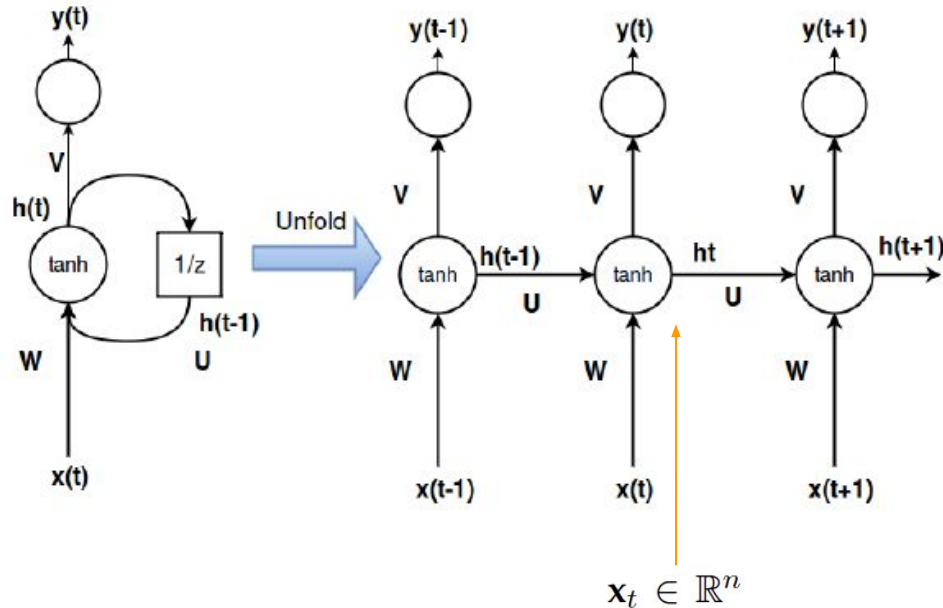
Hence we have two data flows: **Forward in layers + time** propagation

**BEWARE:** We have extra depth now! Every time-step is an extra level of depth (as a deeper stack of layers in a feed-forward fashion!)



# Recurrent Neural Network

Hence we have two data flows: **Forward in layers + time** propagation

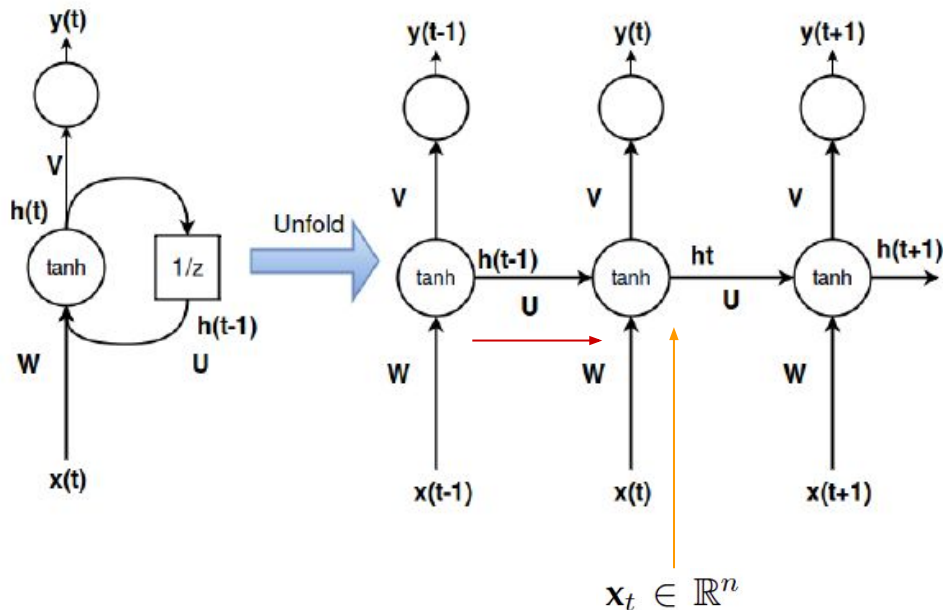


$$h_t = g(W \cdot x_t + U \cdot h_{t-1} + b_h)$$

# Recurrent Neural Network

Hence we have two data flows: **Forward in layers + time** propagation

- Last time-step includes the context of our decisions recursively



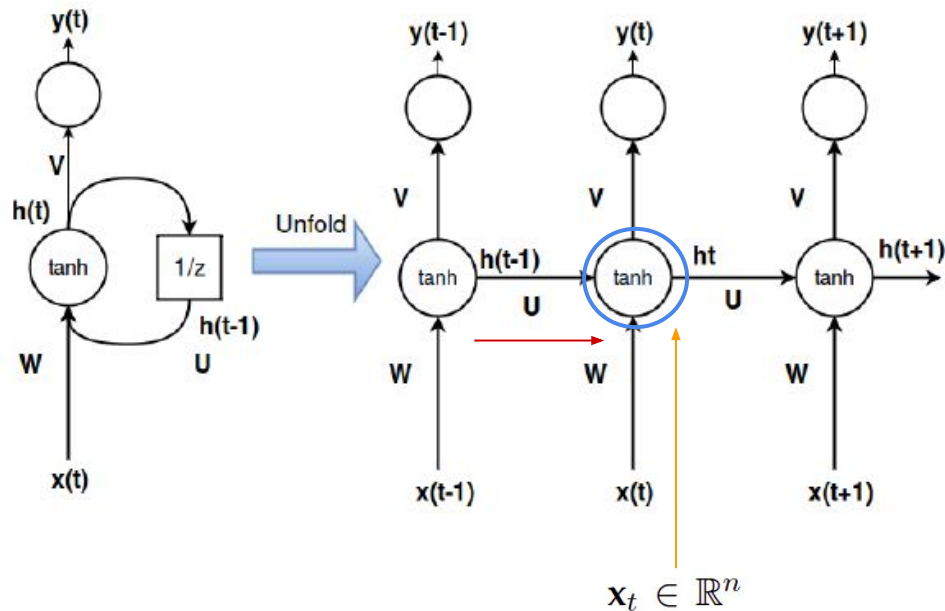
$$h_t = g(W \cdot x_t + U \cdot h_{t-1} + b_h)$$



# Recurrent Neural Network

Hence we have two data flows: **Forward in layers + time** propagation

- Last time-step includes the context of our decisions recursively



$$h_t = g(W \cdot x_t + U \cdot h_{t-1} + b_h)$$

# Recurrent Neural Network

**Back Propagation Through Time (BPTT):** The training method has to take into account the time operations → a cost function  $E$  is defined to train our RNN, and in this case the total error at the output of the network is the sum of the errors at each time-step:


$$E(\mathbf{y}, \hat{\mathbf{y}}) = \sum_{t=1}^T E_t(\mathbf{y}_t, \hat{\mathbf{y}}_t)$$

$$\frac{\partial E}{\partial \mathbf{W}} = \sum_{t=0}^{T-1} \frac{\partial E_t}{\partial \mathbf{W}}$$

**T: max amount of time-steps to do back-prop.** In Keras this is specified when defining the “input shape” to the RNN layer, by means of:  
(batch size, sequence length (T), input\_dim)

Input shape

3D tensor with shape (nb\_samples, timesteps, input\_dim).



Example back-prop in time with 3 time-steps


# Recurrent Neural Network

## Main problems:

- **Long-term memory** (remembering quite far time-steps) **vanishes quickly** because of the recursive operation with **U**

$$\mathbf{h}_t = g(\mathbf{W} \cdot \mathbf{x}_t + \mathbf{U} \cdot g(\cdots g(\mathbf{W} \cdot \mathbf{x}_{t-T} + \mathbf{U} \cdot \mathbf{h}_{t-T} + \mathbf{b}_h) \cdots) + \mathbf{b}_h)$$

- **During training gradients explode/vanish easily because of depth-in-time** → Exploding/Vanishing gradients!




Example back-prop in time with 3 time-steps

# Long Short-Term Memory (LSTM)

1744

Sepp Hochreiter and Jürgen Schmidhuber



Hochreiter, Sepp, and Jürgen Schmidhuber. ["Long short-term memory."](#) Neural computation 9, no. 8 (1997): 1735-1780.

# Long Short-Term Memory (LSTM)

The New York Times, [“When A.I. Matures, It May Call Jürgen Schmidhuber ‘Dad’”](#)  
(November 2016)




# Long Short-Term Memory (LSTM)



Jürgen  
Schmidhuber @  
NIPS 2016  
Barcelona



# Long Short-Term Memory (LSTM)



1991: SEPP HOCHREITER'S ANALYSIS OF THE  
FUNDAMENTAL DEEP LEARNING PROBLEM

$$\left\| \frac{\partial e(t-q)}{\partial e(t)} \right\| = \left\| \prod_{m=1}^q WF'(Net(t-m)) \right\|$$
$$\leq (\|W\| \max_{Net} \{ \|F'(Net)\| \})^q$$

COMPARE: HOCHREITER & BENGIO & FRASCONI & SCHMIDHUBER, 2001

<http://www.idsia.ch/~juergen/fundamentaldeeplearningproblem.html>

[Jürgen Schmidhuber](#) @  
NIPS 2016 Barcelona

# Long Short-Term Memory (LSTM)



[Jürgen Schmidhuber](#) @  
NIPS 2016 Barcelona



# Gating method




Solutions:

1. Change the way in which past information is kept → create the notion of **cell state**, a memory unit that keeps long-term information in a safer way by protecting it from recursive operations
2. Make every RNN unit able to **forget whatever may not be useful anymore** by clearing that info from the cell state (optimized clearing mechanism)
3. Make every RNN unit able to decide whether **the current time-step information matters or not**, to accept or discard (optimized reading mechanism)
4. Make every RNN unit able to **output the decisions whenever it is ready to do so** (optimized output mechanism)

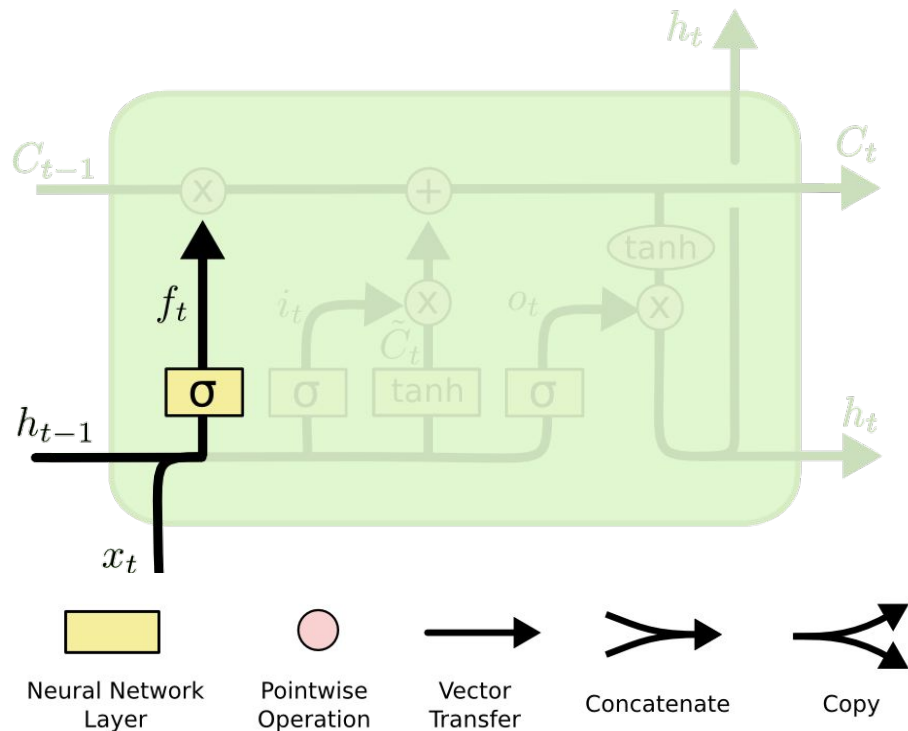
# Long Short-Term Memory (LSTM)

Three **gates** are governed by *sigmoid* units (btw [0,1]) define the control of in & out information..



# Long Short-Term Memory (LSTM)

Make every RNN unit able to **forget whatever may not be useful anymore** by clearing that info from the cell state (optimized clearing mechanism)



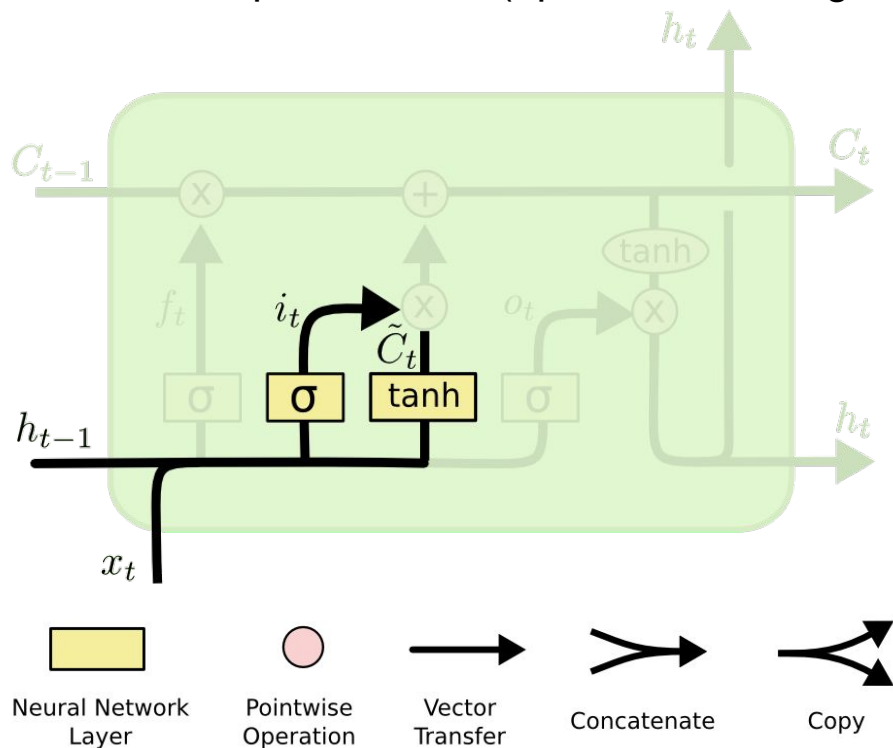
**Forget Gate:**

$$f_t = \sigma \left( W_f \cdot \underbrace{[h_{t-1}, x_t]}_{\text{Concatenate}} + b_f \right)$$

Concatenate

# Long Short-Term Memory (LSTM)

Make every RNN unit able to decide whether **the current time-step information matters or not**, to accept or discard (optimized reading mechanism)



## Input Gate Layer


$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

## New contribution to cell state

$$\tilde{C}_t = \underbrace{\tanh(W_C \cdot [h_{t-1}, x_t] + b_C)}_{\text{Classic neuron}}$$

# Long Short-Term Memory (LSTM)

Make every RNN unit able to decide whether **the current time-step information matters or not**, to accept or discard (optimized reading mechanism)

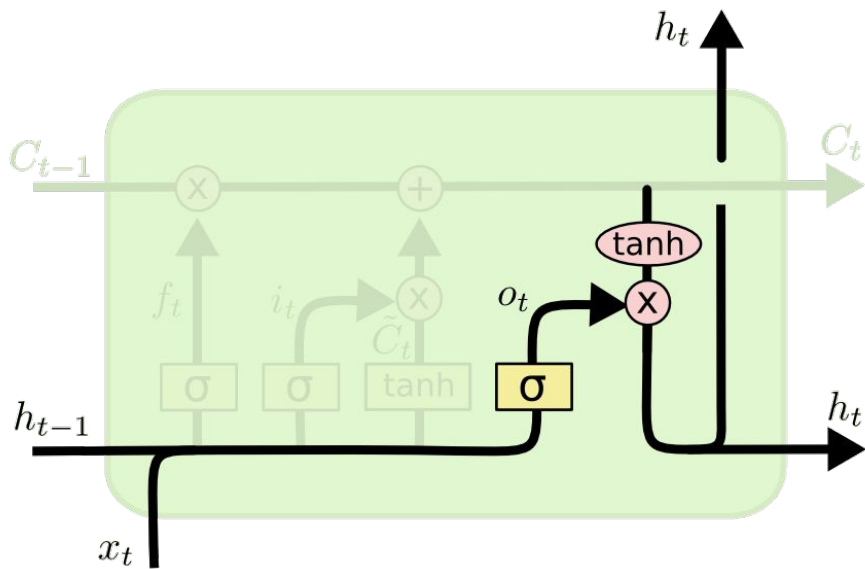


**Update Cell State (memory):**

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

# Long Short-Term Memory (LSTM)

Make every RNN unit able to **output the decisions whenever it is ready to do so** (optimized output mechanism)




## Output Gate Layer

$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

## Output to next layer

$$h_t = o_t * \tanh (C_t)$$


# Long Short-Term Memory (LSTM)



# Long Short Term Memory (LSTM) cell

An LSTM cell is defined by two groups of neurons plus the cell state (memory unit):

1. **Gates**
2. **Activation units**
3. **Cell state**



$$\mathbf{i}_t = \sigma(\mathbf{W}_i \mathbf{x}_t + \mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{b}_i)$$

$$\hat{\mathbf{C}}_t = \tanh(\mathbf{W}_c \mathbf{x}_t + \mathbf{U}_c \mathbf{h}_{t-1} + \mathbf{b}_c)$$


$$\mathbf{f}_t = \sigma(\mathbf{W}_f \mathbf{x}_t + \mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{b}_f)$$

$$\mathbf{C}_t = \mathbf{i}_t \odot \hat{\mathbf{C}}_t + \mathbf{f}_t \odot \mathbf{C}_{t-1}$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o \mathbf{x}_t + \mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{b}_o)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{C}_t)$$

Computation  
Flow






# Long Short Term Memory (LSTM) cell

An LSTM cell is defined by two groups of neurons plus the cell state (memory unit):

1. **Gates**
2. **Activation units**
3. **Cell state**


$$N_{params}^i = 4 \times (N_{inputs}^i \times N_{units}^i + N_{units}^i \times N_{units}^i + N_{units}^i)$$

3 gates + input activation



# Gated Recurrent Unit (GRU)

Similar performance as LSTM with less computation.



$$u_i = \sigma \left( W^{(u)} x_i + U^{(u)} h_{i-1} + b^{(u)} \right) \quad (1)$$

$$r_i = \sigma \left( W^{(r)} x_i + U^{(r)} h_{i-1} + b^{(r)} \right) \quad (2)$$

$$\tilde{h}_i = \tanh \left( W x_i + r_i \circ U h_{i-1} + b^{(h)} \right) \quad (3)$$


$$h_i = u_i \circ \tilde{h}_i + (1 - u_i) \circ h_{i-1} \quad (4)$$

$$N_{params}^i = 3 \times (N_{inputs}^i \times N_{units}^i + N_{units}^i \times N_{units}^i + N_{units}^i)$$

Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. ["Learning phrase representations using RNN encoder-decoder for statistical machine translation."](#) AMNLP 2014.

# Gated Recurrent Unit (GRU)

Similar performance as LSTM with less computation.



$$u_i = \sigma \left( W^{(u)} x_i + U^{(u)} h_{i-1} + b^{(u)} \right) \quad (1)$$

$$r_i = \sigma \left( W^{(r)} x_i + U^{(r)} h_{i-1} + b^{(r)} \right) \quad (2)$$


$$\tilde{h}_i = \tanh \left( W x_i + r_i \circ U h_{i-1} + b^{(h)} \right) \quad (3)$$

$$h_i = u_i \circ \tilde{h}_i + (1 - u_i) \circ h_{i-1} \quad (4)$$

$$N_{params}^i = 3 \times (N_{inputs}^i \times N_{units}^i + N_{units}^i \times N_{units}^i + N_{units}^i)$$

Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. ["Learning phrase representations using RNN encoder-decoder for statistical machine translation."](#) AMNLP 2014.

# Gated Recurrent Unit (GRU)




$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# Thanks ! Q&A ?



Follow me at



[/ProfessorXavi](#)



[@DocXavi](#)



UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

Department of Signal Theory  
and Communications

*Image Processing Group*

<https://imatge.upc.edu/web/people/xavier-giro>